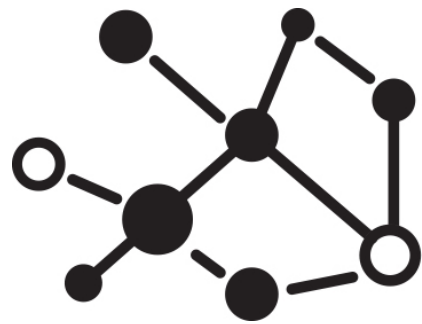


Stream Processing with MicroProfile and Apache Kafka



MicroProfile



Apache Kafka

Background: Motivation for MicroService and Kafka



Agenda

- **MicroProfile**
- WildFly Swarm
- Apache Kafka
- Integrating MicroProfile and Kafka
- CDI Extension for Apache Kafka
- Outlook

Enterprise Java Standards History



Java
Community
Process



RED HAT JBOSS
ENTERPRISE
APPLICATION PLATFORM 7

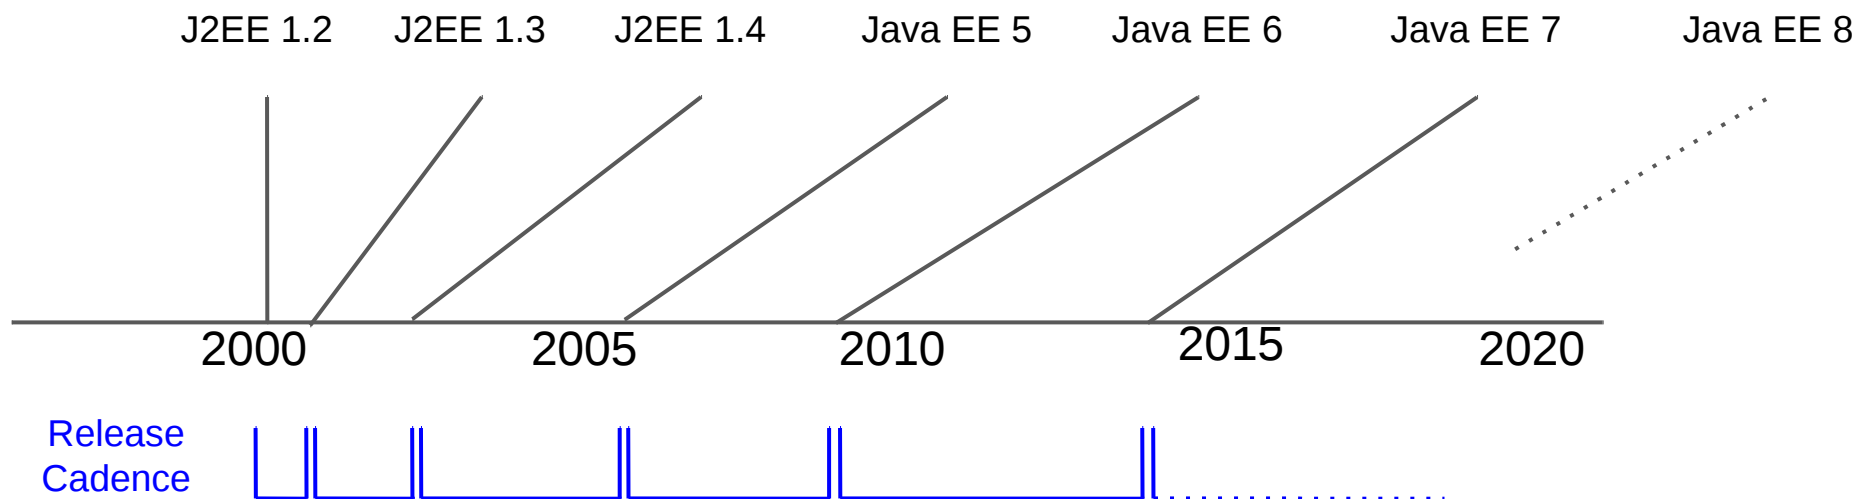
ORACLE
FUSION MIDDLEWARE
WEBLOGIC SERVER



Cosminexus

FUJITSU

TmaxSoft



MicroProfile Background

Began as a collection of independent discussions

Many innovative “microservices” efforts in existing Java EE projects

[WildFly Swarm](#)

[WebSphere Liberty](#)

[Payara](#)

[TomEE](#)

Projects already leveraging both Java EE and non-Java EE technologies

Creating new features/capabilities to address microservices architectures

Quickly realized there is common ground

Java EE technologies are already being used for microservices,

but we can do better

Matthias Wessendorf – Red Hat | matzew AT redhat DOT com | @mwessendorf

MicroProfile Release Philosophy

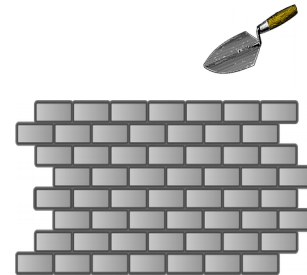
Release 1.0



Rapidly iterate
and innovate



Build
consensus



Standardize



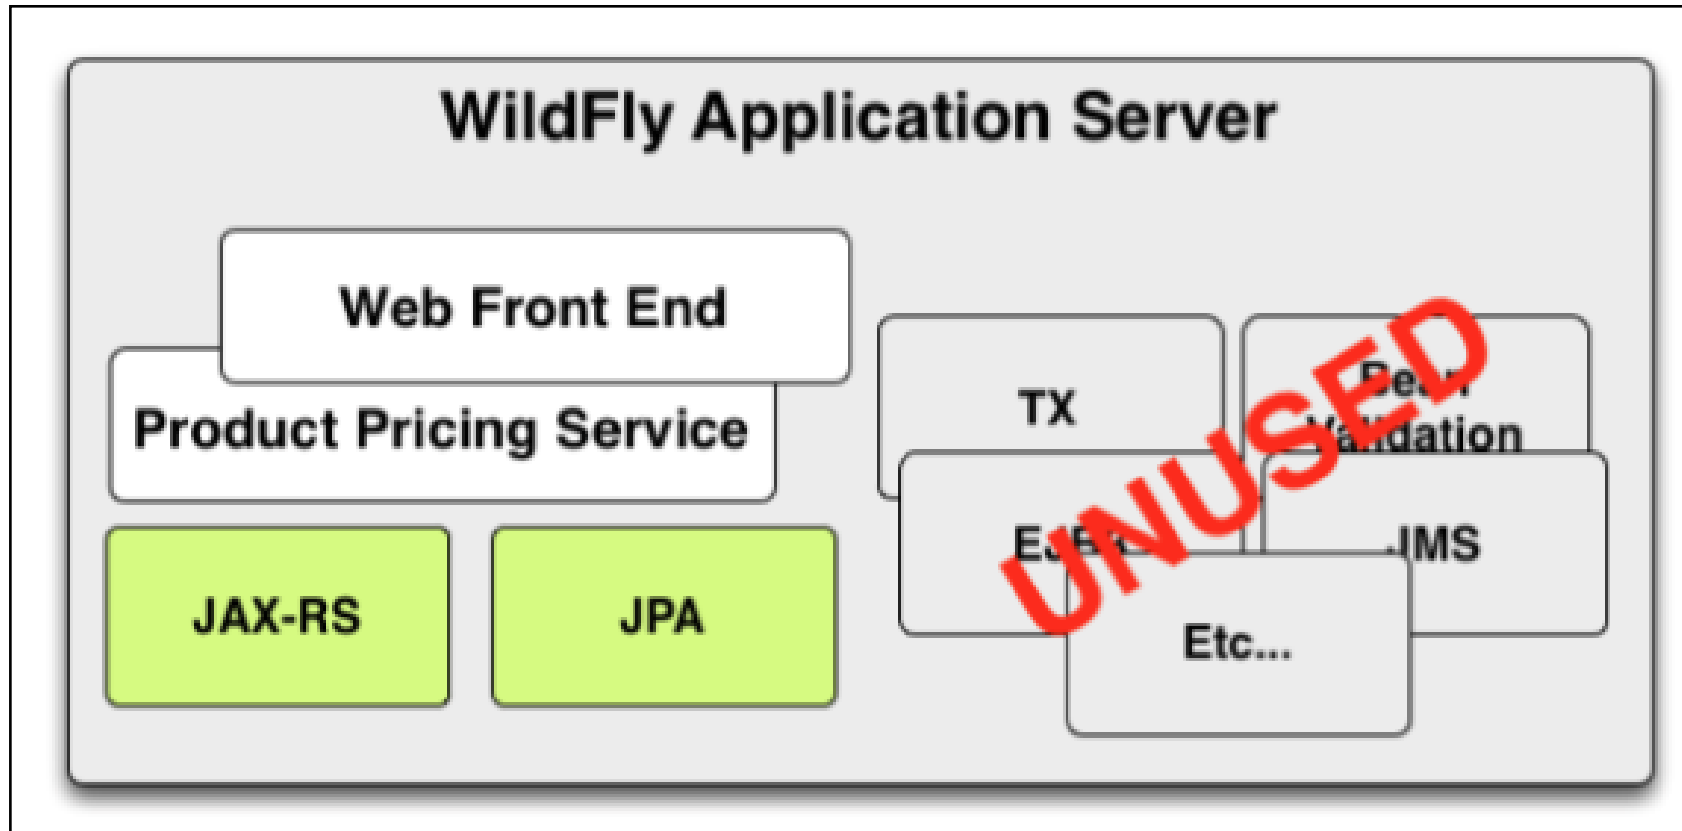
Sept 2016

Agenda



- MicroProfile
- **WildFly Swarm**
- Apache Kafka
- Integrating MicroProfile and Kafka
- CDI Extension for Apache Kafka
- Outlook

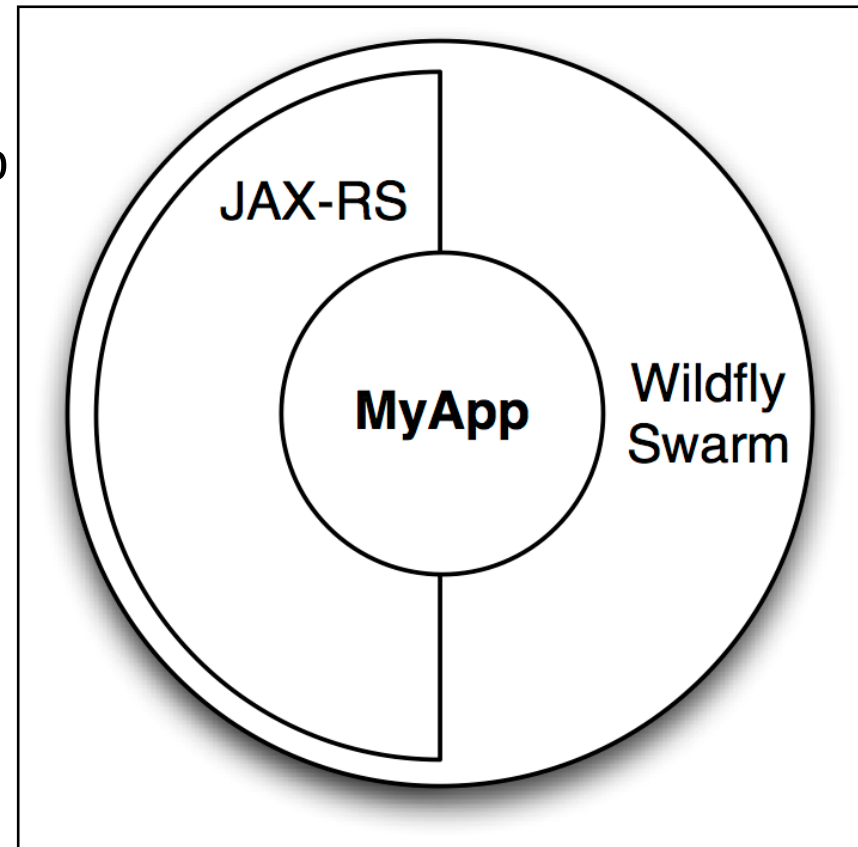
Just Enough App Server



- Use the API's you want
- Include the capabilities you need
- Wrap it up for deployment

Uberjar

- A single .jar file containing your application,
- the portions of WildFly required to support it,
- an internal Maven repository of dependencies,
- plus a shim to bootstrap it all



Fractions

- A well-defined collection of application capabilities.
 - May map directly to a WildFly subsystem,
 - or bring in external capabilities such as Netflix Ribbon.

What Fractions do

- Enable WildFly subsystems (JAX-RS, Infinispan)
- Integrate additional system capabilities (Topology)
- Provide deployments (ribbon-webapp, jolokia)
- Alter deployments (keycloak)

DEMO

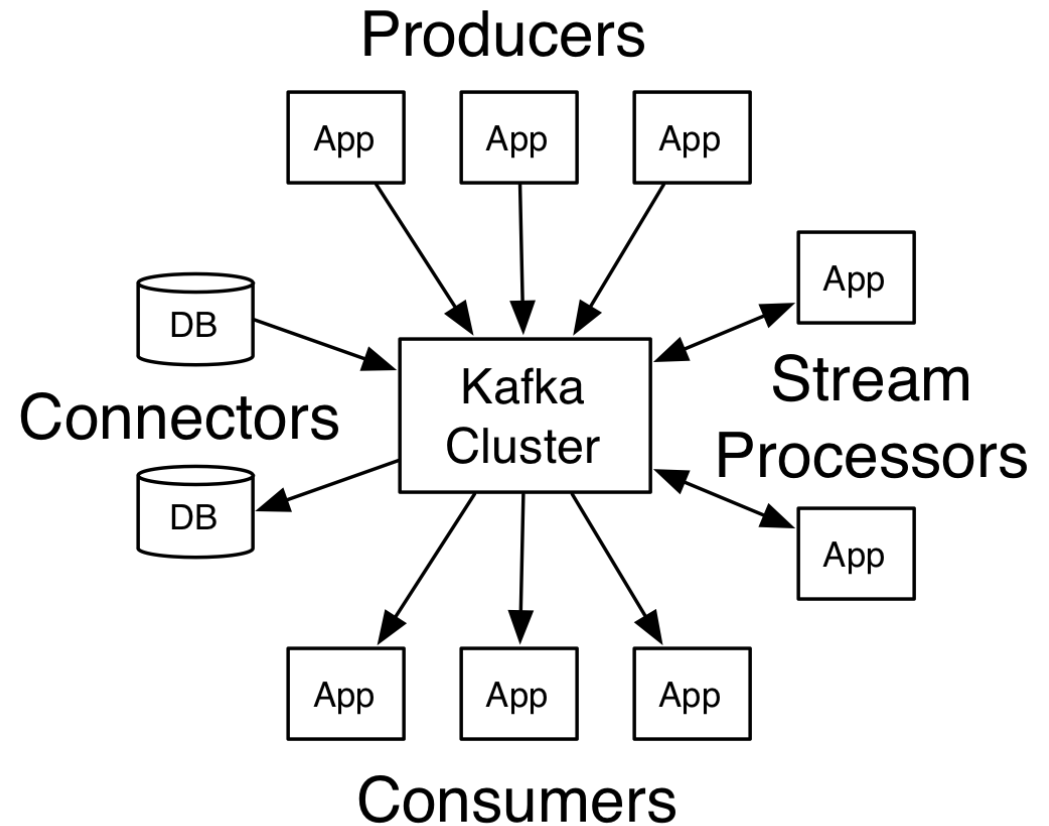
... MicroProfile Server App using Swarm!

Agenda

- MicroProfile
- WildFly Swarm
- **Apache Kafka**
- Integrating MicroProfile and Kafka
- CDI Extension for Apache Kafka
- Outlook

Apache Kafka

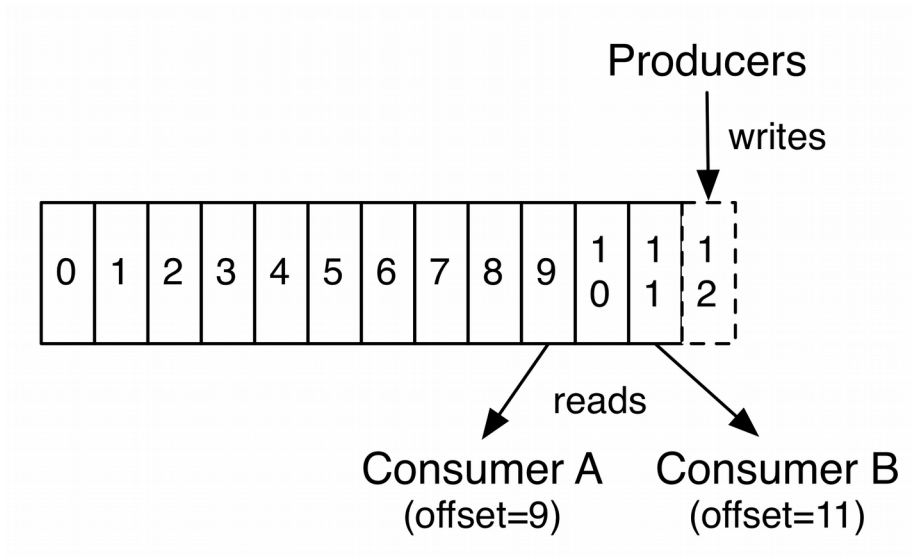
- like messaging system, but different
 - “distributed commit log”
- Clustering is CORE...
- Durability & Ordering Guarantees
- Typical Use-Cases
 - ETL / **C**hange **D**ata **C**apture
 - <http://debezium.io> (CDC)
 - Data Pipeline: Kafka as the HUB for other systems
 - User activity tracking/reporting
 - analytics....



DEMO

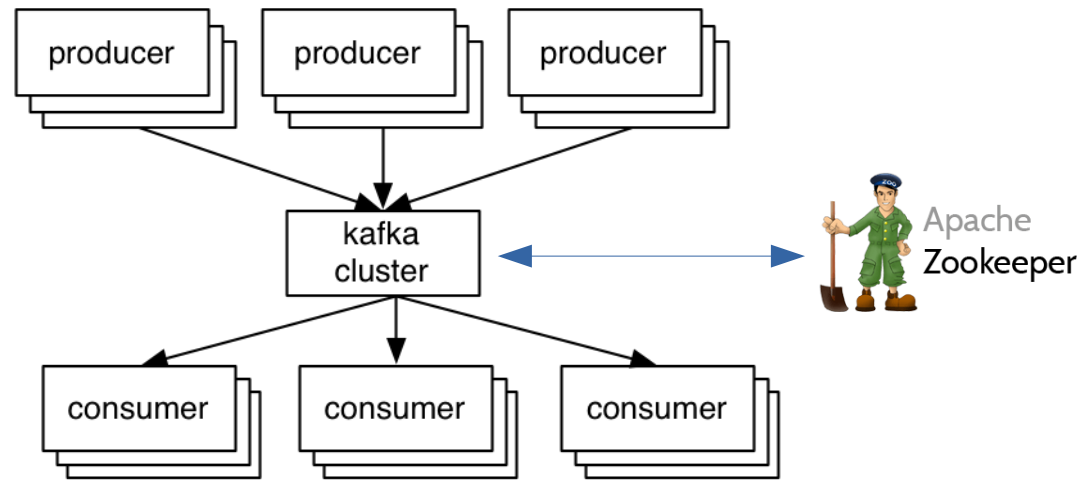
From WebSocket to Apache Kafka

Records (or Messages)



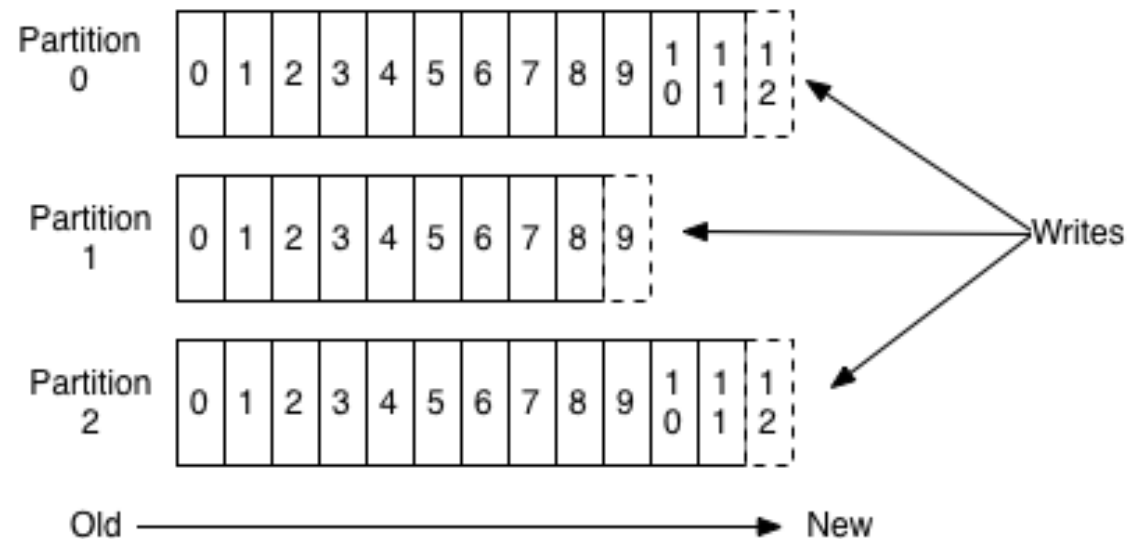
- Byte Array
 - Key/Value pairs
- Immutable
- Records (or messages, or events) are being appended
- Persisted to disk

Producers and Consumers



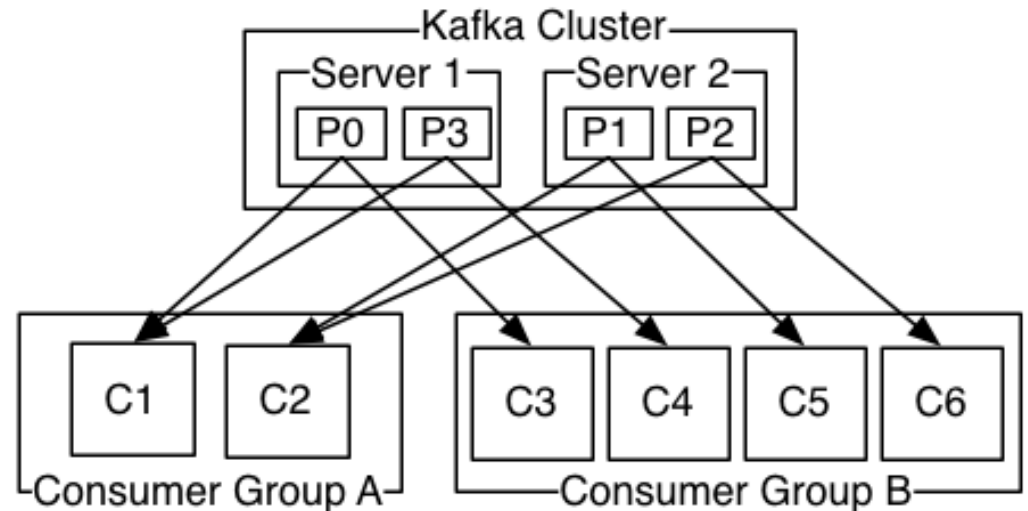
- n nodes/brokers → Kafka cluster (clients connect to bootstrap servers)
 - Apache Zookeeper
- Producer sends message to a broker
- Consumer is connected to a broker, and *polls* message from a broker
- Leader/Follower architecture...

Topics, Partitions and Offsets



- Topic is containing 1 or more partitions
 - Guaranteed ordering (“only” on a Partition of a Topic)
- Replication of the partitions (Leader/Follower)
 - Partitioning-Factor (per Topic) is configured when setting up a Topic
- Offset: unique sequential ID per TopicPartition
- Consumer keeps track of offset
 - Reply or handling consumers with different speed! :-)

Consumer Groups



- Logical grouping of some Kafka consumers
 - groups receive msg from Topic: **AT_LEAST_ONCE**
 - individual consumer: assigned to partition(s) of the cluster
- Separate *scaling* for each consumer group (listening on same Topic)
 - Example:
 - Group A: expensive/non-time-sensitive → scale down....
 - Group B: realtime processing / time-sensitive → scale up

DEMO

WebSocket demo: behind the scenes...


Some details on Apache Kafka's Java API (0.10.2.0)

Agenda

- MicroProfile
- WildFly Swarm
- Apache Kafka
- **Integrating MicroProfile and Kafka**
- CDI Extension for Apache Kafka
- Outlook

Integration: Kafka and Microprofile

- Kafka's Java library is easy to integrate
- Wiring of Producers and Consumers with CDI
- Contexts and Dependency Injection (CDI) for the Java EE platform
 - Contexts: The ability to bind the lifecycle and interactions of stateful components to well-defined but extensible lifecycle contexts
 - Dependency injection: The ability to inject components into an application in a typesafe way, including the ability to choose at deployment time which implementation of a particular interface to inject
- CDI is intended to be a foundation for frameworks, extensions and integration with other technologies!

```
public class ProcessorBean {  
  
     @Inject  
    private KafkaProducer<String, String> producer;  
  
    public void publishEventRecord(final String payload) {  
        producer.send(new ProducerRecord<String, String>(topic: "some_topic", payload));  
    }  
}
```

```
@Produces
public KafkaProducer createProducer() {

    final Properties properties = new Properties();

    properties.put(BOOTSTRAP_SERVERS_CONFIG, "172.17.0.3:9092");
    properties.put(KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class);
    properties.put(VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class);

    return new KafkaProducer<>(properties);
}
```

```
@ApplicationScoped
public class ConsumerRegistrationFactory {

    Logger logger = Logger.getLogger(ConsumerRegistrationFactory.class.getName());

    // This will add a new thread to our pool, to subscribe to our Observable
    @Resource(name = "DefaultManagedExecutorService")
    private ManagedExecutorService executor;

    public void init(@Observes @Initialized(ApplicationScoped.class) Object init) {
        logger.severe("msq: Setup my consumer");

        // create (all) consumer(s)
        MyKafkaConsumer myKafkaConsumer = new MyKafkaConsumer();

        // submit to managed executor service
        executor.submit(myKafkaConsumer);
    }
}
```


Agenda

- MicroProfile
- WildFly Swarm
- Apache Kafka
- Integrating MicroProfile and Kafka
- **CDI Extension for Apache Kafka**
- Outlook

CDI *portable* extensions for Apache Kafka



- CDI is intended to be a foundation for frameworks, extensions and integration with other technologies!
 - Customize the platform for individual needs!
- Removes boilerplate code, makes Kafka usage really easy!
- CDI extension requires 3 “things”
 - beans.xml (*optional since CDI 1.1*)
 - services file
 - Implementation class: POJO observing the CDI lifecycle events
- CDI: A ***great!*** way for extending the standardized platform!
 - Hence it was critical for MicroProfile too!

Meet kafka-cdi

... A simple CDI extension for Apache Kafka

<https://github.com/matzew/kafka-cdi>

Agenda

- MicroProfile
- WildFly Swarm
- Apache Kafka
- Integrating MicroProfile and Kafka
- CDI Extension and Swarm Fraction
- **Outlook**

AeroGear UPS

POC:

- Swarm based JAX-RS endpoint for Push
- Kafka as the event stream
- Consumer to process Push Metrics (e.g. from Apple)

There is more.... much more...! Believe me, it's true!

- KStream API
 - New API, build on-top of Kafka's Java client
 - Functional programming to filter/map/reduce streams
 - No need for complexer frameworks like Spark or Flink
- Vert.x
 - Nice and simple wrapper around Kafka's Java client
- Debezium platform for CDC
 - contains `KafkaCluster` class for testing!, or demos :-)
- Future options:
 - More CDI / Swarm enhancements (e.g. JCA, Swarm Fraction)

THANKS!

Questions ?

Beer !

Cocktails!



Slides and (some) demos:

<https://github.com/matzew/kafka-microprofile>